

Introduction to Artificial Intelligence

Homework 7 Resolution by Dino Meng [SM3201466]

Q1. MDPs: Reward Shaping

(a)

In this exercise, let us denote each tiles of the board as a couple (\bullet, \bullet) :




Figure 1: Board Notation in this exercise

For instance, the tile with -1 will be identified as $(B, 5)$ in this exercise.

Let us proceed the resolution.

Observe that for each non-empty tile exiting is never a convenient choice, as it nets them a discounted reward of 0 and they can always get something “better” by heading towards a non-empty tile with positive reward and exit there. So we will exclude from said tiles the choice of *exitting*.

Now we will deduct the optimal choice for each tile by calculating possible rewards for each choice taken.

$(A, 1)$: The only choice is to go down as observed before, so we will mark \downarrow

$(B, 1)$: We can either *exit*, *go down* or *go right*.

- Going down makes no sense as from there the only sensible choice would be to exit (otherwise I would return to the same tile, “wasting” reward as they are “weakened” from the discount $\gamma = 1/2$)
- Immediately exiting yields a reward of 1
- Going right, right, down and exiting yields a reward of $0 + (1/2)0 + (1/4)0 + (1/8)4 = 1/2$

So the best choice would be to immediately exit, so we mark **EXIT**

$(C, 1)$: We can either go up or immediately exit. * Exit: Yields a reward of 1 * Up: The next node is $(B, 1)$ where I should exit. The reward yielded is 1, which then gets saturated to $1/2$ due to the discount factor

So we mark **EXIT** as it’s the best choice

$(B, 2)$: We can either go left or right.

- Left: We can then immediately exit, yielding the discounted sum of rewards $1/2$
- Right: We can go down and exit at $(C, 3)$, yielding $0 + 1/2(0) + 1/4(4) = 1$

The best choice is then to go right, so we mark \rightarrow

$(B, 3)$: Similar reasoning to $(B, 2)$, except that it's even more evident that heading towards $(C, 3)$ is better. So we mark \downarrow

$(C, 3)$: The best course of action would be to exit, as it would yield the maximum reward 4 in the whole scenario. So EXIT

$(C, 4)$: We can either immediately exit or go left or go right.

- Right: Similarly to case $(B, 1)$, going right makes no sense as I get no reward or even a negative value
- Left: I can exit on $(C, 3)$, yielding the discounted reward $(1/2)4 = 2$
- Exit: Yields 1

The best choice is therefore \leftarrow

$(C, 5)$: Similar analysis to $(C, 4)$. So the best choice is \leftarrow

$(B, 5)$: In this case, exiting has a negative outcome so we must move "somewhere" to avoid that. In this case \downarrow as it's the only choice.

To recap,




Figure 2: solution, grid form

(b). Firstly, we can calculate each value F explicitly as the state-action space is sufficiently small.

s	a	s'	F
A	\rightarrow	B	10
A	EXIT	/	0
B	\leftarrow	A	0
B	\rightarrow	C	10

s	a	s'	F
B	EXIT	/	0
C	←	B	0
C	EXIT	/	0

Now, to compute the optimal policies we proceed state by state.

A: Of course →, as it's the only action that yields a reward

B: We can exclude exiting as not only it does not produce a reward, it leads us to a terminal state without a chance of getting further rewards.

- Right: We immediately receive a reward of 10
- Left: We can go right again, receiving the discounted reward $(1/2)10 = 5$

So → is the optimal choice

C: Exiting leads us to a “dead end” without any rewards. However, if we were to go left and right again, we would receive the discounted reward 5. In fact, if we were to repeat this pattern forever we would converge towards “something positive”, since that it's bounded by a convergent series:

$$0 \leq \sum_{i=1,3,5,\dots,(2n+1),\dots} \frac{1}{2^i} 10 \leq \sum_{i=1} \frac{1}{2^i} 10 = \sum_{i=0} \frac{1}{2^i} 10 - 10 = \frac{10}{1-0.5} - 10 = 20 - 10 = 10$$

So the best choice is ← as it would at least potentially yield some sort of reward




Figure 3: recap b

(c). 0, because the agent never exits thus never getting any reward from the original environment

(d). As observed previously, the problem lies in the fact that the agent might “be encouraged” to go “back and forth” forever to get an increasing reward. So one can “discourage” the agent from doing that by setting x a sufficiently negative value such that the cost of “going further from C” outweighs the reward of “approaching towards C”. For example, $x = -11$ causes the desired situation.

Q2. Micro-Blackjack

(a). We will fill a table of triples (s, a, s') and their associated probabilities and rewards

s	a	s'	P	R
0	Draw	2	1/3	0
	Draw	3	1/3	0
	Draw	4	1/3	0
	Stop	End	1	0
2	Draw	4	1/3	0
	Draw	5	1/3	0
	Draw	End	0	0

s	a	s'	P	R
	Stop	End	1	2
3	Draw	5	1/3	0
	Draw	End	2/3	0
	Stop	End	1	3
4	Draw	End	1	0
	Stop	End	1	4
5	Draw	End	1	0
	Stop	End	1	5

N.B. If a field is empty, then it's the same as row above (column s)

To fill out the table, I have proceeded with the following “algorithm”:

- For each state 0,2,4,5:
 - If draw, consider each possibility and sum it with the original state; if it's below 6, then let s' the sum, otherwise let it be the end with reward 0
 - If stop, let it be end with reward of the number of the state

(b).

Step 0: All of the values are 0, by definition

Step 1: Note that in all states it is never convenient to draw, as we will always get reward 0 by Bellman's equation for fixed policy “draw” (since that V_0 is initialized to zeroes). So $V_1(s) = s$.

Step 2:

- 0: If I drew, then the expected reward would be $2/3 + 3/3 + 4/3 = 9/3 = 3$. Otherwise, by staying the reward would have been 0. So by taking the max. the value is 3
- 2: If I drew, then the expected reward would be $4/3 + 5/3 = 9/3 = 3$, which is better than stopping as the reward would have been 2. So by max. the value is 3
- 3: Here drawing is not convenient anymore, as the expected reward would be $5/3$ which is lower than 3, obtained by stopping. So the value is 3
- 4: Remains the same, drawing is not convenient at all
- 5: Same as 4

Step 3: Note that the values for the states 4, 5 are stabilized and as a consequence so are 2, 3 as they depend on these values. The only one which can change is 0

- 0: By drawing, the expected reward is $3/3 + 3/3 + 4/3 = 10/3 = 1 + 1/3$. Therefore, the value is $1 + 1/3$.

Step 4: Same as step 3, the values have stabilized

(c). By repeating the reasoning done in (b), we conclude that:

- Drawing is convenient for 0,2
- Stopping is convenient for 3,4,5

States	0	2	3	4	5
V_0	0	0	0	0	0
V_1	0	2	3	4	5
V_2	3	3	3	4	5
V_3	$3+1/3$	3	3	4	5
V_4	$3+1/3$	3	3	4	5

Figure 4: point b

States	0	2	3	4	5
π^*	Draw	Draw	Stop	Stop	Stop

Figure 5: point c